Learning the Rhythm: Why RoPE Breaks Beyond the Training Horizon, and How Learnable Frequencies Fix It

Liu Yuncong Zhang Yihang Ma Xiaoyu

FT5011 Deep Learning for Finance
National University of Singapore

11 November 2026

Contents

1. Problem & Background

2. Method

3. Experiments & Mechanism

4. Application to Finance

5. Limitations & Takeaways

Problem & Background

Problem: one model, many horizons

Finance needs forecasts at several horizons at once.

- Risk: Value-at-Risk at **1 day** and **10 days**; volatility daily $\rightarrow$ weekly $\rightarrow$ monthly (HAR).
- Energy trading: bid for **tomorrow**, hedge for **next week**.
- Usual practice: *one model per horizon* — costly to train and maintain.

Goal

Train once on a short horizon, then forecast **any** horizon:

$$\underbrace{x_{t-335}, \ldots, x_t}_{L=336 \text{ past steps}} \longmapsto \underbrace{\hat{x}_{t+1}, \ldots, \hat{x}_{t+H'}}_{\text{trained with } H=96,\ \text{asked for } H' \leq 720}$$

A **step** = one observation. On hourly prices, 96 steps = **4 days**, 720 steps = **30 days**.

The problem: silent failure beyond the training horizon

What happens. Trained for 96 steps (4 days), asked for 720 (30 days): the forecast breaks down beyond the training horizon.

Why it matters.

- **Silent**: fine at the training horizon, no warning beyond it.
- **Common**: RoPE is standard in LLMs and time-series models.
- **Costly**: the alternative is one model per horizon.

6.4×

error at 720 steps vs. a model trained for 720 (Electricity)

Why it happens (pages 12, 18–19)

RoPE's **fixed periods** do not match the **daily cycle**: each day the same hour gets a slightly different position code, so far ahead the model meets codes it never saw.

Where we started: the proposal

Proposal idea.

- HAR fixes horizons by hand (day / week / month).
- Give each head a **learnable time scale** on its rotary encoding; let the model *learn* which scales matter.

What we actually found.

- At the training horizon the scales **barely move** ($\pm 10\%$); accuracy unchanged.
- Heads do **not** specialise to different scales.
- Beyond the training horizon they make a **huge** difference.

This talk

Why does RoPE break beyond the training horizon, why do 4 parameters per layer fix it, and what does it mean for finance?

Rotary Position Embedding (RoPE)

Definition (Su et al., 2021, *RoFormer*)

For each attention head of dimension d_h (even), split into $m = d_h/2$ two-dimensional blocks. At position ℓ, RoPE applies a rotation with frequency ω_i:

$$\mathrm{RoPE}(x, \ell) = \begin{bmatrix} x_{2i}\cos(\omega_i \ell) - x_{2i+1}\sin(\omega_i \ell) \\ x_{2i}\sin(\omega_i \ell) + x_{2i+1}\cos(\omega_i \ell) \end{bmatrix}, \qquad \omega_i = \theta^{-2i/d_h}.$$

Key Property

The dot product between a query–key pair after RoPE depends only on their **relative offset**:

$$\langle \mathrm{RoPE}(q, \ell_q),\, \mathrm{RoPE}(k, \ell_k) \rangle = \langle q,\, R(\omega_i(\ell_k - \ell_q))\, k \rangle.$$

Hence attention sees **relative distances** only — what matters for time series is the lag, not the calendar index.

RoPE Mechanism — Frequency Interpretation

Expand one 2D block. Let $q^{(i)} = [q_{i,1}, q_{i,2}]^\top$, $k^{(i)} = [k_{i,1}, k_{i,2}]^\top$:

$$q^{(i)\top} R(\omega_i \Delta) k^{(i)} = (q_{i,1}k_{i,1} + q_{i,2}k_{i,2})\cos(\omega_i\Delta) + (q_{i,2}k_{i,1} - q_{i,1}k_{i,2})\sin(\omega_i\Delta).$$

Define coefficients.

$$a_i = \langle q^{(i)}, k^{(i)} \rangle, \quad b_i = q_{i,2}k_{i,1} - q_{i,1}k_{i,2}.$$

Trigonometric form

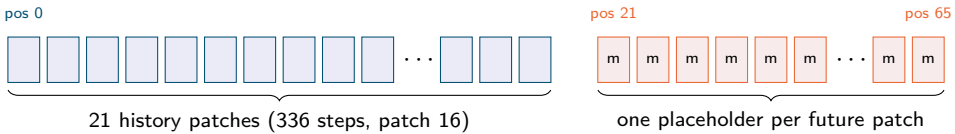
$$S(\Delta) = \sum_{i=0}^{m-1} [a_i \cos(\omega_i \Delta) + b_i \sin(\omega_i \Delta)].$$

Interpretation. Each block acts as a **sinusoidal filter**: large $\omega_i \to$ short-term/high-frequency; small $\omega_i \to$ long-term/low-frequency.

Head-wise scaling (our proposal). $\omega_i^{(h)} = \omega_i \, e^{a_h}$ was meant to let heads specialise on distinct horizons. We find instead that the learned scale aligns the **fastest** block ($\omega_0 = 1$ rad/token, period $2\pi \approx 6.28$ tokens) with the daily cycle.

Method

Model: placeholders for the future



- 2-layer, 4-head Transformer encoder; $d = 64$. Each history patch of 16 steps is one token.
- All placeholders are the **same** learnable vector `m`: **only RoPE tells them apart**.
- Mask: history sees history; a placeholder sees history and itself $\Rightarrow$ the forecast for step h does not depend on H, so we can ask for **any** horizon after training.
- Trained with $H = 96$ (6 placeholders); queried with $H = 720$ (45 placeholders).

Our method: learnable frequency scales

Per-head learnable scale (the proposal's idea)

$$\hat{\omega}_{\ell,h,j} = \omega_j \cdot e^{a_{\ell,h}}, \qquad a_{\ell,h} \text{ learned, initialised at } 0 \ (= \text{standard RoPE})$$

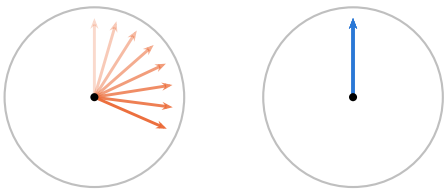
Only **4 extra parameters per layer** (one per head).

Ablation variants (to find out *why* it works):

- *Grouped*: separate scales for the fast and the slow frequency pairs.
- *Shared*: one scale for all heads of a layer.
- *Fixed*: scales set by hand and frozen (spread across heads, or aligned to the day).

Second remedy — multi-horizon training: each batch uses $H \in \{96, 192, 336, 720\}$, so one model also sees distant placeholders.

Intuition: RoPE as a clock



Standard RoPE ($\omega_0 = 1$) Aligned ($\omega_0 = 2\pi/6$)

fastest hand at midnight on days 0–7 (darker = later)

- ETTm1: 15-min data, patch 16 $\Rightarrow$ **1 day = 6 tokens**.
- Standard RoPE turns once per **6.28** tokens: it falls behind by 0.28 tokens a day.
- Inside the training window the drift is small; 7 days ahead the hand points somewhere **never seen in training**.

Alignment

If $\hat{\omega}_0 = 2\pi/T$ (T = tokens per day), position $t + T$ looks exactly like t: the far future looks like the past.

Experiments & Mechanism

Experimental setup

Data (7 datasets, every variable its own series)

- Benchmarks: Electricity, ETTh1, ETTm1, Weather.
- **Financial**: day-ahead electricity prices Nord Pool (EPF-NP) & PJM (EPF-PJM); daily exchange rates.
- Chronological splits; z-score on training rows only; RevIN for exchange rates.

Metric: MSE on the normalised scale.

Baselines & why

- *Fixed RoPE*: same model without scales $\Rightarrow$ isolates the mechanism.
- *Linear* $\mathbb{R}^{336} \to \mathbb{R}^{H}$: as strong as Transformers here (DLinear).
- *Trained directly at 720*: the natural alternative to extrapolation.

Protocol & significance

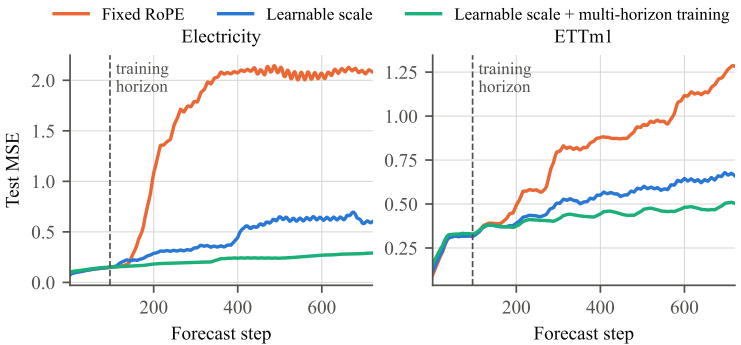
- Same backbone, paired seeds, base and checkpoint chosen on validation only.
- 3–5 seeds + **Diebold–Mariano** test (Newey–West, overlapping windows).

Result 1: at the training horizon

Test MSE, $H = 96$	Linear	Transformer, fixed RoPE	Transformer, learnable scale
Electricity	0.140	0.139	0.139
ETTh1	**0.377**	0.417	0.414
ETTm1	**0.304**	0.311	0.317
Weather	0.174	**0.153**	**0.153**
EPF-NP	0.472	0.482	**0.452**
EPF-PJM	**0.163**	0.174	0.167
Exchange	**0.089**	0.103	0.099

- The Transformer is close to published results (PatchTST: Electricity 0.129, Weather 0.149).
- Linear models remain strong (as reported by DLinear).
- Learnable scales neither help nor hurt here — the effect lies **beyond** the horizon.

Result 2: fixed RoPE collapses beyond it



Test MSE per forecast step; trained at $H = 96$ (dashed line).

6.4×

error of fixed RoPE at 720 steps vs. a model trained at 720 (Electricity)

38–66%

of that error removed by learnable scales

6 / 6

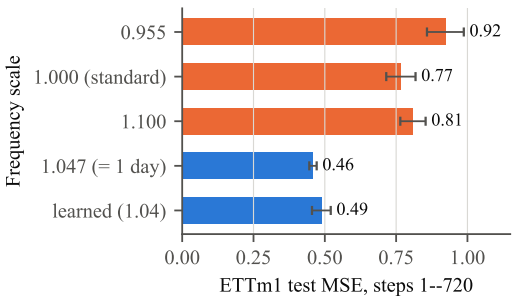
runs significant (Diebold–Mariano) on Electricity, ETTh1, ETTm1

Result 2: all datasets at 720 steps

	Trained at 96, asked for 720		DM	Multi-horizon	Trained at 720	
Test MSE, steps 1–720	fixed	learnable	+ / −	learnable	Transformer	Linear
Electricity	1.347±0.087	0.464±0.107	6 / 0	0.216±0.003	0.210±0.002	0.204±0.000
ETTh1	0.972±0.071	0.473±0.013	6 / 0	0.541±0.014	0.632±0.036	0.493±0.004
ETTm1	0.764±0.068	0.474±0.019	6 / 0	0.421±0.011	0.435±0.013	0.430±0.003
Weather	0.392±0.009	0.396±0.006	2 / 1	0.326±0.009	0.336±0.004	0.325±0.001
EPF-NP	1.072±0.061	0.896±0.039	5 / 0	1.123±0.008	1.196±0.156	0.863±0.003
EPF-PJM	0.443±0.036	0.273±0.034	6 / 0	0.216±0.002	0.213±0.003	0.221±0.000
Exchange	1.078±0.100	1.024±0.036	3 / 3	1.074±0.087	1.097±0.214	0.890±0.003

- Strong, significant gains on seasonal data: Electricity, ETTh1, ETTm1 and **both electricity markets**.
- No gain on Weather and Exchange rates (DM: mixed signs) — explained on the next slides.
- DM column: of 6 runs (3 seeds $\times$ 2 bases), how many are significantly better / worse ($p < 0.05$).

Why: alignment to the daily cycle

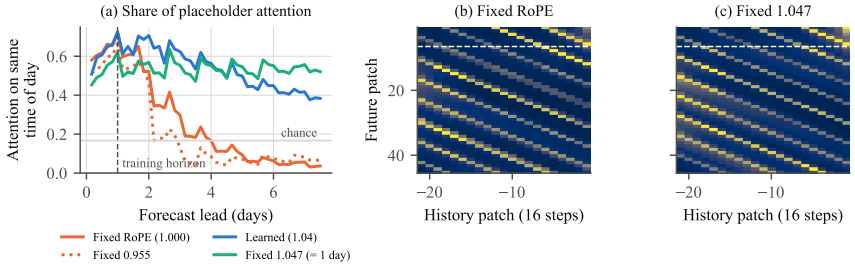


- ETTm1: 1 day = 6 tokens $\Rightarrow$ aligned scale $2\pi/6 = \mathbf{1.047}$.
- Learned scales average **1.04** — in both layers.
- **Fixed at 1.047, no learning**: as good as learning.
- 5% too slow (0.955) or fast (1.10): collapses like standard RoPE.
- DM test 1.047 vs 0.955: $p < 10^{-15}$ in 6/6 runs.

Mechanism

Learning works because it **tunes RoPE's fastest rotation to the daily cycle**.

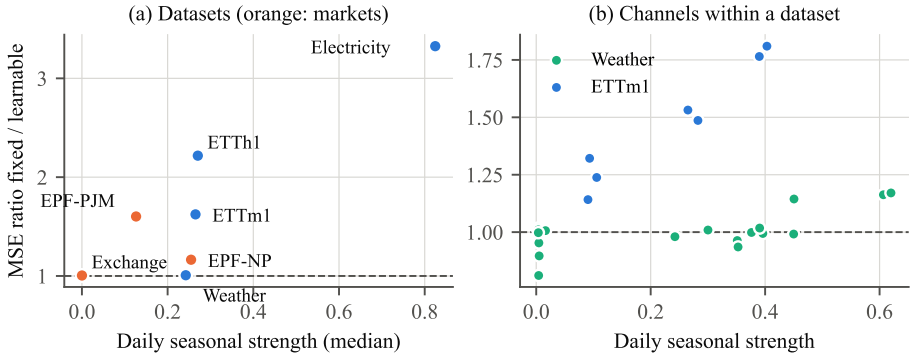
Attention: what the placeholders look at



ETTm1, trained at $H = 96$, queried at 720; last layer, mean of 3 seeds; dashed line = training horizon.

- Future patches attend to history at the **same time of day** (stripes every 6 patches): a "same time yesterday" lookup.
- **Fixed RoPE loses it** after $\sim$2 days ($0.64 \to 0.05$, below chance); **aligned 1.047 keeps it** (≈ 0.53); learned 1.04 decays slowly.

Gain grows with the daily cycle



- **Datasets**: strongest cycle (Electricity) gains most; no cycle (Exchange) gains nothing.
- **Channels**: in Weather only the 3 solar-radiation channels gain — so Weather as a whole shows none.

Remedy 2: multi-horizon training

MSE, 720 steps	Learnable @96	Multi-H (1 model)	Trained @720
Electricity	0.464	**0.216**	0.210
ETTm1	0.474	**0.421**	0.435
Weather	0.396	**0.326**	0.336
EPF-PJM	0.273	**0.216**	0.213
ETTh1	**0.473**	0.541	0.632
EPF-NP	**0.896**	1.123	1.196

- One model trained on mixed horizons is within 3% of, or better than, the model trained at 720 — on **all 7 datasets**.
- Small cost at $H = 96$ (Electricity 0.135 vs 0.128).
- With it, fixed and learnable RoPE perform alike: distant positions are no longer unseen.
- When long-horizon training **overfits** (small ETTh1, noisy EPF-NP), learnable scales + extrapolation are best.

Application to Finance

Application to finance

Electricity markets (daily cycle)

Learnable scales cut 720-step error significantly:
EPF-PJM $0.44 \rightarrow \mathbf{0.27}$ (6/6 runs),
EPF-NP $1.07 \rightarrow \mathbf{0.90}$ (5/6 runs).

Exchange rates (random walk)

No cycle to align $\Rightarrow$ no consistent effect (3 better, 3 worse); linear model best; RevIN essential.

Practical rules

1. For **seasonal** financial series — intraday volume & volatility (daily U-shape), power & gas prices — use learnable or calendar-aligned RoPE before forecasting beyond the training horizon.
2. To cover several risk horizons (1- & 10-day VaR, HAR-style volatility) with **one model**, train on mixed horizons.
3. For near-random-walk prices, do not expect position encodings to help.

Limitations & Takeaways

Limitations & future work

What did not work / limitations

- Heads did not specialise to distinct time scales (the proposal's hypothesis).
- Grouping frequencies (the starting repository's hypothesis) gave no extra gain.
- Small model (2 layers, $d = 64$), 3–5 seeds (partly offset by DM tests).
- With point tokens instead of patches the effect is much noisier; linear models stay strong at the training horizon.

Future work

- Initialise frequencies from the data's dominant period (FFT) instead of learning them.
- Align several frequencies to several cycles (day *and* week).
- Test on larger Transformers and time-series foundation models that use RoPE.

Takeaways

1. The problem

RoPE forecasters trained at one horizon **collapse** beyond it: up to $6.4\times$ the error of a directly trained model.

2. The cause

RoPE's rotation is **out of step with the daily cycle**; the far future looks unlike anything seen in training.

3. The fix

4 learnable scales per layer **align** it (38–66% less error, significant); mixed-horizon training lets **one model** match per-horizon models.

For finance

Works on seasonal markets (electricity prices); not on random-walk prices (exchange rates).

Thank you!

Questions?

Backup: What Does *Not* Matter

Frequency scales	Test MSE
ETTh1, steps 1–336, 5 seeds	
Fixed RoPE (scale 1)	0.676±0.022
Learnable, per head	0.460±0.006
Learnable, grouped	0.460±0.006
Learnable, shared by heads	0.458±0.011
Fixed, mild spread (0.9–1.1)	0.547±0.038
Fixed, wide spread (0.5–2.0)	0.716±0.064
ETTm1, steps 1–720, 3 seeds × 2 bases	
Fixed RoPE (scale 1)	0.767±0.051
Learnable, per head	0.488±0.033
Fixed 1.047 (= 1 day)	0.459±0.013
Fixed 0.955 (misaligned)	0.923±0.065
Fixed 1.10 (misaligned)	0.809±0.044

- **Head diversity**: one scale shared by all heads = one per head. Hand-set spreads help little or hurt.
  ⇒ heads do *not* specialise to time scales.
- **Grouping**: the fast group learns the same value as the per-head scale; the slow group drifts with no effect.
  ⇒ only the **fast** frequencies matter.
- **Gradient clipping**: inactive after warm-up — ruled out.

Backup: Transformer Basics

Architecture. Standard encoder–decoder with positional encodings (Vaswani et al., 2017).

Tokenization. History as tokens $\{x_1, \ldots, x_L\}$ (points or patches).

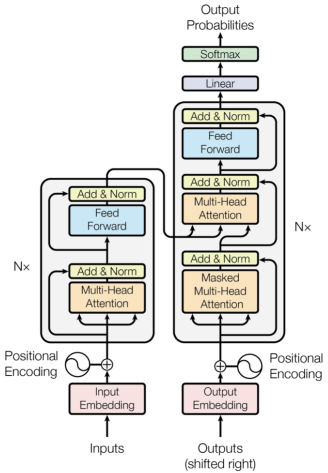
Projections.

$Q = XW_Q, \quad K = XW_K, \quad V = XW_V \quad (\in \mathbb{R}^{L \times d})$

Scaled dot-product attention (key equation)

$$\mathrm{Attn}(Q, K, V) = \mathrm{Softmax}\left(\frac{QK^\top}{\sqrt{d}}\right) V$$

Multi-head attention allows different heads to capture diverse temporal patterns. Stacking layers builds contextual representations for sequence prediction.



Backup: Positional Encoding

Why do we need position encodings?

Self-attention treats the input as an unordered set of tokens — without additional signals, it cannot tell which time step comes first.

- To restore temporal order, each token x_ℓ is added with a **positional encoding**:

$$\tilde{x}_\ell = x_\ell + PE(\ell).$$

- The encoding $PE(\ell)$ can be:
  - **Sinusoidal:** fixed periodic functions $\sin(\omega\ell), \cos(\omega\ell)$ (Vaswani et al., 2017).
  - **Learnable:** parameters trained with the model (BERT-style).
- This addition injects absolute position information into the input before attention.

RoPE's idea

RoPE **replaces this absolute positional addition step** with a **rotation-based operation** inside the attention itself — making the attention score aware of relative offsets instead of absolute indices.

Backup: Why Relative Positions

- In NLP, **absolute positions** work because word order carries meaning.
- In time-series forecasting, especially in **financial data**, we care more about **relative offsets**:
  - Forecasting at $t{+}1$ should be invariant to where the window starts — the market pattern "5 days before a crash" should look the same in 2018 or 2024.
  - Temporal similarity depends on **distance** Δ, not on the absolute calendar index ℓ.
  - Financial series are often non-stationary — absolute time (e.g., day 2000 vs. day 4000) has little predictive meaning.

Motivation for RoPE

$$\underbrace{\begin{aligned} q_m &= f_q(x_m, m), \\ k_n &= f_k(x_n, n) \end{aligned}}_{\text{absolute position encoding}} \quad \Longrightarrow \quad \underbrace{\langle f_q(x_m, m), f_k(x_n, n) \rangle = g(x_m, x_n, m-n)}_{\text{relative position encoding target}}$$

Backup: From Rotation to Relative Encoding

Setup. Let d_h be even, $m = d_h/2$. Split $q, k \in \mathbb{R}^{d_h}$ into 2D blocks:

$$q = [q^{(0)}; \ldots; q^{(m-1)}], \quad k = [k^{(0)}; \ldots; k^{(m-1)}], \quad q^{(i)}, k^{(i)} \in \mathbb{R}^2.$$

RoPE rotation.

$$R(\omega_i \ell) = \begin{bmatrix} \cos(\omega_i \ell) & -\sin(\omega_i \ell) \\ \sin(\omega_i \ell) & \cos(\omega_i \ell) \end{bmatrix}, \quad q'^{(i)} = R(\omega_i \ell_q) q^{(i)}, \; k'^{(i)} = R(\omega_i \ell_k) k^{(i)}.$$

Dot-product score:

$$S = \frac{(q')^\top k'}{\sqrt{d_h}} = \frac{1}{\sqrt{d_h}} \sum_{i=0}^{m-1} q'^{(i)\top} k'^{(i)}.$$

Goal

Show S can be rewritten as a function of the **relative offset** $\Delta = \ell_k - \ell_q$.

Backup: Trigonometric Foundation of RoPE

Rotation matrix: $R(\theta) = \begin{bmatrix} \cos\theta & -\sin\theta \\ \sin\theta & \cos\theta \end{bmatrix}$.

Compute:

$$R(\alpha)^\top R(\beta) = \begin{bmatrix} \cos\alpha & \sin\alpha \\ -\sin\alpha & \cos\alpha \end{bmatrix} \begin{bmatrix} \cos\beta & -\sin\beta \\ \sin\beta & \cos\beta \end{bmatrix}.$$

Multiplication result:

$$= \begin{bmatrix} \cos\alpha\cos\beta + \sin\alpha\sin\beta & -\cos\alpha\sin\beta + \sin\alpha\cos\beta \\ -\sin\alpha\cos\beta + \cos\alpha\sin\beta & \sin\alpha\sin\beta + \cos\alpha\cos\beta \end{bmatrix}.$$

Apply trig identities:

$$\begin{cases} \cos(\beta-\alpha) = \cos\beta\cos\alpha + \sin\beta\sin\alpha, \\ \sin(\beta-\alpha) = \sin\beta\cos\alpha - \cos\beta\sin\alpha, \end{cases}$$

Hence:

$$R(\alpha)^\top R(\beta) = \begin{bmatrix} \cos(\beta-\alpha) & -\sin(\beta-\alpha) \\ \sin(\beta-\alpha) & \cos(\beta-\alpha) \end{bmatrix} = R(\beta-\alpha).$$

Composing two rotations equals a single rotation by the **angle difference** $(\beta - \alpha)$ — the key property enabling RoPE's relative encoding.

Backup: Relative Dependence Emerges

Key trigonometric identity.

$$R(\alpha)^\top R(\beta) = R(\beta - \alpha) \quad \Longleftrightarrow \quad \begin{cases} \cos(\beta-\alpha) = \cos\beta\cos\alpha + \sin\beta\sin\alpha, \\ \sin(\beta-\alpha) = \sin\beta\cos\alpha - \cos\beta\sin\alpha. \end{cases}$$

Apply to RoPE rotations.

$$q'^{(i)\top} k'^{(i)} = (R(\omega_i \ell_q) q^{(i)})^\top (R(\omega_i \ell_k) k^{(i)}) = q^{(i)\top} R(\omega_i(\ell_k - \ell_q)) k^{(i)}.$$

Therefore

$$\boxed{S(\Delta) = \frac{1}{\sqrt{d_h}} \sum_{i=0}^{m-1} q^{(i)\top} R(\omega_i \Delta) k^{(i)}}, \qquad \Delta = \ell_k - \ell_q.$$

Here the trigonometric identity merges the query and key rotations into a single relative rotation $R(\omega_i\Delta)$, so each 2D block interacts only through the relative offset Δ.

Backup: replication & hyper-parameters

- Code: `github.com/ycsama0703/ft5011-ropeformer` + `experiments/` (commands in `experiments/README.md`).
- Data: Time-Series-Library (HuggingFace, CC BY 4.0); EPF (Zenodo 4624805); exchange rates (Time-Series-Library).
- $L = 336$, patch 16, $d = 64$, 4 heads, 2 layers, FFN 128, dropout 0.1, batch 256.
- AdamW, lr 10^{-3} (no weight decay on scales), 300 warm-up steps + cosine, clip 1.0, bf16.
- Budgets: 2000 steps (ETTh1, EPF, exchange), 4000 (ETTm1), 6000 (Weather), 8000 (Electricity).
- RoPE base $\in \{10^3, 10^4\}$ and best checkpoint chosen on validation; test used once.
- One run: 20 s – 3 min on one H100 (memory capped at 8%).